

VO BA HOANG MINH

AI ENGINEER | EVALUATION, AGENTS & RETRIEVAL

Ho Chi Minh City, Vietnam | Remote, worldwide | UTC+7

minhhoang250803@gmail.com | +84 902 391 936

github.com/anhminhzui-dev | [Portfolio \(anhminhzui-dev.github.io\)](https://portfolio.anhminhzui-dev.github.io) | [Research \(doi.org/10.5281/zenodo.22675789\)](https://doi.org/10.5281/zenodo.22675789)

Founder of Gnomon, delivering paid writing-assessment work for teachers and clients. Engineering experience spans data science, retrieval, multimodal training and evaluation.

EXPERIENCE

Gnomon | Founder & AI Systems Engineer | 2024 - Present

- Delivered paid writing-assessment outputs for teachers and clients; owned product design, issue triage, corrective fixes and final review.
- Built exemplar retrieval for Gnomon and a separate LangGraph pipeline that retrieves compliance rules with SQLite FTS5.
- Engineered selection across 101,175 records and 47 source labels; prepared 10,888 training and 578 validation records using 44 data-quality checks.
- Led multimodal supervised fine-tuning with completion-only loss and chunked vocabulary projection; a separate 8B run processed 9,325 examples, including 2,134 with images.
- Designed agent scheduling with capacity and write-scope guardrails; directed React/TypeScript and FastAPI/SQLite delivery.

FPT Software | Data Science | 2025 - 2026

- Consolidated customer and business data, resolving missing records, inconsistencies and duplicates; analysed sales and operational data for trends, seasonality and anomalies.
- Compared forecasting models against baselines to recommend model selection.

FPT Education | Lecturer, AI for VFX | 2025 - 2026

- Delivered paid instruction on applying AI to visual-effects workflows.

CONSULTING EXPERIENCE

Forward-deployed engineering | VFX and education

- Delivered AI, data and pipeline work with VFX and education client teams, alongside full-time roles.

Government digitisation | OCR books to e-books

- Worked with contractor teams on Vietnam's OCR book-to-e-book digitisation programme.

PROJECTS

failclosed-eval | Input validation | [View](#)

- Built fail-closed input validation for images, leakage and duplicates: 375 of 500 ELLIPSE-based payloads admitted, 125 malformed inputs refused.

policy-deck | Command classification | [View](#)

- Benchmarked 11 command-classification rules: 8 false positives/0 false negatives on 1,260 fitted cases; 2/0 on 190 separate cases. Classifier, not an execution sandbox.

relay-gate | Completion-claim detection | [View](#)

- Built a detector for false agent completion claims; benchmarked on MAST and AgentRewardBench, with 81 tests.

mcp-trajectory-judge | Tool-trajectory evaluation | [View](#)

- Replays synthetic MCP-style tool trajectories in a deterministic sandbox to check actions against completion claims.

PUBLICATIONS & OPEN SOURCE

Constructing Language Judgments: Rubric Distillation and Evidence-Grounded Reasoning | [View](#)

- Authored a technical report on rubric compilation, evidence verification and resource-aware multimodal training.

EleutherAI lm-evaluation-harness | Submitted PR #4115 | [View](#)

- Added isolated DummyLM tests for construction, generation/error paths and chat templates; pull request open, not yet merged.

SKILLS

Evaluation: structured-output validation, pytest, GitHub Actions CI.

Retrieval and agents: LangGraph, theta-kNN exemplar retrieval, SQLite FTS5.

Machine Learning: QLoRA, multimodal SFT, data-quality and leakage checks.

Languages, tools and serving: Python, PyTorch, PEFT/LoRA, TypeScript, React, FastAPI, SQLite.

Spoken languages: Vietnamese: native | English: professional proficiency | Japanese: intermediate.

EDUCATION

FPT University | Bachelor's degree.