

VO BA HOANG MINH

AI ENGINEER | APPLIED LLM SYSTEMS & EVALUATION

Ho Chi Minh City, Vietnam | Remote from HCMC / HCMC onsite or hybrid | UTC+7
minhhoang250803@gmail.com | +84 902 391 936

[GitHub](#) | [Portfolio - all projects & research](#) | [Research \(DOI\)](#)

AI engineer combining data-science experience at FPT Software with paid assessment and workflow delivery through Gnomon. Designs grounded LLM systems, training pipelines and agent orchestration; teaches AI applications at FPT Education.

EXPERIENCE

Gnomon | Founder & AI Systems Engineer | 2024-present

- Delivered paid writing-assessment outputs and AI/data workflows to teachers and clients; led product design, issue identification, corrective work and output review.
- Built the retrieval layer (nearest-neighbour exemplar retrieval with a kNN-derived range constraint) inside the production assessment engine, and a LangGraph-orchestrated pipeline with a full-text (BM25-family) retrieval step, 4 of 4 tests passing.
- Engineered data selection across 101,175 records and 47 source labels; prepared 10,888 training and 578 validation records. The data system includes 44 registered checks for source, task, image, duplication and leakage risks.
- Led multimodal SFT with completion-only weighted loss, chunked vocabulary projection and accumulation-consistent normalization; a separate 8B run used 9,325 examples, 2,134 with images.
- Engineered rubric distillation and structured reasoning supervision: compiled item banks, source-span verification, relation-level credit accounting and candidate-bound corrective examples for judgment construction.
- Designed dependency-aware agent scheduling, capacity and write-scope checks; directed React/TypeScript and FastAPI/SQLite workflows from assignment through teacher review and learner release.

FPT Software | Data Science | 2025-early 2026

- Prepared and consolidated historical customer and business data for analysis and predictive-model evaluation, addressing missing records, inconsistencies and duplication.
- Conducted exploratory analysis of sales, customer behaviour and operational data to identify trends, seasonality, relationships and unusual changes.
- Compared forecasting models against baselines to recommend model selection; evaluated predictive and text-model outputs across business forecasting, risk/anomaly detection and language tasks.

FPT Education | Lecturer, AI for VFX | 2025, 2026

- Delivered paid instruction on applying AI to visual-effects workflows.

CONTRACT & CONSULTING WORK | alongside full-time roles

Forward-deployed engineering on client contracts | VFX and education industries

- Embedded with client teams to deliver AI, data and pipeline work on contract, across the VFX and education industries, alongside full-time roles.

Government digitisation programme | OCR conversion of printed books to e-books

- Worked with contractor teams on Vietnam's book-to-e-book OCR digitisation programme to preserve printed knowledge.

SELECTED ENGINEERING PROJECTS | 2026

LLM Evaluation & Data Quality | failclosed-eval | [View](#)

- Tested missing images, label leakage and duplicates: 375 valid inputs admitted; 125 deliberately malformed inputs rejected in a controlled 500-payload ELLIPSE-based test.

Agent Command-Policy Evaluation | policy-deck | [View](#)

- Evaluated 11 command-classification rules on 1,260 fitted synthetic cases and 190 separate, partly rule-aware cases; separate-deck result: two false positives, zero false negatives.

COMPETITIONS | SEPTEMBER 2026

Voice-agent hackathon (lablab.ai) | ARC White-Box Estimation 2026 (Alcrowd) | Kaggle S6E9 | OpenCV | [View](#)

- Submitted voice-agent entry with public code (uh-huh, 2 companion tools); scored ARC white-box leaderboard submission (arc-whest); Kaggle S6E9 public score 0.94151; OpenCV entry registered.

RESEARCH & OPEN-SOURCE CONTRIBUTION

Constructing Language Judgments: Rubric Distillation and Evidence-Grounded Reasoning | [View](#)

- Authored a systems report connecting rubric compilation, evidence-verifying inference harnesses, verifier-filtered supervision and resource-aware multimodal adaptation.

EleutherAI lm-evaluation-harness | Submitted PR #4115 | [View](#)

- Added hermetic DummyLM tests for construction, generation/error paths and chat-template behavior; pull request open, not yet merged.

SKILLS, EDUCATION & LANGUAGES

Python, SQL/SQLite, TypeScript, FastAPI, REST APIs, React, pytest, GitHub Actions; PyTorch, PEFT/QLoRA, multimodal SFT, LangGraph, retrieval, data quality, structured-output validation, evaluation design.

Google Cloud / Vertex AI: model APIs, retrieval, fine-tuning, endpoints and prompt/quality/latency/cost evaluation.

Google Developer Program member; prototyped website/API protection for an AI project. FPT University | Bachelor's degree. Vietnamese: native | English: professional proficiency | Japanese: intermediate.