

VO BA HOANG MINH

AI ENGINEER | EVALUATION, AGENTS & RETRIEVAL

Ho Chi Minh City, Vietnam | Remote, worldwide | UTC+7

minhhoang250803@gmail.com | +84 902 391 936

github.com/anhminhzui-dev | [Portfolio \(anhminhzui-dev.github.io\)](https://Portfolio(anhminhzui-dev.github.io)) | [Research \(doi.org/10.5281/zenodo.22675789\)](https://Research(doi.org/10.5281/zenodo.22675789))

Founder of Gnomon, delivering paid writing-assessment work for teachers and clients. Engineering experience spans data science, retrieval, multimodal training and evaluation.

EXPERIENCE

Gnomon | Founder & AI Systems Engineer | 2024 - Present

- Delivered paid teacher and client assessments; owned product design, triage, fixes and final review.
- Designed the facts-not-grades architecture: the model emits evidence grounded in source spans, a deterministic scorer maps facts to the result, and an abstain layer refuses on thin evidence.
- Built the fail-closed evaluation harness: 2,232 criterion units across 629 essays, 1,866 admitted and 366 refused before scoring; LLM-as-judge runs on a frozen, hash-checked evidence window.
- Built KD-tree exemplar retrieval for Gnomon and a separate LangGraph pipeline that retrieves compliance rules with SQLite FTS5.
- Engineered selection across 101,175 records and 47 source labels; prepared 10,888 training and 578 validation records using 44 data-quality checks.
- Led multimodal supervised fine-tuning: QLoRA on a 27B open-weight model (4-bit NF4) with completion-only loss and chunked vocabulary projection; a separate 8B run processed 9,325 examples, 2,134 with images.
- Directed a two-provider coding-agent fleet under policy-as-code guardrails (resource-aware admission, spend/irreversibility classifier scored on a hash-pinned 1,219-case deck, output-hash delivery checks); shipped the React/TypeScript and FastAPI/SQLite product: 43 router modules, 79 migrations, TOTP MFA, 1,117 backend tests.
- Presentation design and slide reconstruction of a ten-slide Gnomon sales deck: action titles, data visualization, rationale log.

FPT Software | Data Science | 2025 - 2026

- Consolidated customer and business data (missing records, inconsistencies, duplicates); analysed sales and operational data for trends, seasonality and anomalies; compared forecasting models to recommend selection.

FPT Education | Lecturer, AI for VFX | 2025 - 2026

- Delivered paid instruction on applying AI to visual-effects workflows.

FORWARD-DEPLOYED ENGINEERING - VFX AND AI | 2023 - 2025

OTSU Labs | Applied AI Operations, VFX Production | 2025

- Ran Applied AI Operations across VFX production projects in partnership with Sparta VFX and Sparx; owned render passes, asset management and look development, delivering shots to studio standard on deadline; converted realistic assets into stylised anime visuals with the 2D team.

FPT, "Fragments of the Deep" | Applied AI Operations Lead | 07/2024 - 2025

- Led the artist team on an Unreal Engine 5 animated short, owning storyboard and camera direction; optimised real-time models for a 30% performance gain.

Dihaan Media | Applied AI Operations | 2023

- Embedded with the client production team on stage and booth environments for Sofitel Saigon Plaza; introduced IES-based lighting and new tools into their workflow.

Contract engagements | VFX, education and government | 2023 - present

- AI, data and pipeline work with VFX and education client teams; contractor work on Vietnam's OCR book-to-e-book digitisation programme.

PROJECTS

failclosed-eval | Input validation | [View](#)

- Built fail-closed input validation for images, leakage and duplicates: 375 of 500 ELLIPSE-based payloads admitted, 125 malformed inputs refused; 239 tests, GitHub CI.

policy-deck | Command classification | [View](#)

- Benchmarked 11 command-classification rules: 8 false positives/0 false negatives on 1,260 fitted cases; 2/0 on 190 separate cases; 455 tests. Classifier, not an execution sandbox.

relay-gate | Completion-claim detection | [View](#)

- Built a detector for false agent completion claims; benchmarked on MAST and AgentRewardBench, with 81 tests.

mcp-trajectory-judge | Tool-trajectory evaluation | [View](#)

- Replays MCP-style tool trajectories in a deterministic sandbox to check actions against completion claims.

Hackathon and competition entries | September 2026 | [View](#)

- Submitted a voice-agent entry with public code (uh-huh, 2 companion tools); scored an ARC White-Box Estimation leaderboard submission (arc-whest); Kaggle S6E9 public score 0.94151; OpenCV entry registered.

PUBLICATIONS & OPEN SOURCE

Constructing Language Judgments: Rubric Distillation and Evidence-Grounded Reasoning | [View](#)

EleutherAI Im-evaluation-harness | Submitted PR #4115 | [View](#)

- Added isolated DummyLM tests for construction, generation/error paths and chat templates.

SKILLS

Evaluation: LLM-as-judge on hashed evidence, structured-output validation, pytest, GitHub Actions CI.

Retrieval and agents: LangGraph, KD-tree theta-kNN exemplar retrieval, SQLite FTS5, policy-as-code agent guardrails.

Machine Learning: Large Language Model (LLM) SFT, QLoRA, data-quality and leakage checks.

Languages, tools and serving: Python, PyTorch, PEFT/LoRA, TypeScript, React, FastAPI, SQLite.

Spoken languages: Vietnamese: native | English: professional proficiency | Japanese: intermediate.

EDUCATION

FPT University | Bachelor's degree.