

VO BA HOANG MINH

AI ENGINEER | EVALUATION, AGENTS & RETRIEVAL

Ho Chi Minh City, Vietnam | Remote, worldwide | UTC+7

minhhoang250803@gmail.com | +84 902 391 936

github.com/anhminhzui-dev | [Portfolio \(anhminhzui-dev.github.io\)](https://Portfolio(anhminhzui-dev.github.io)) | [Research \(doi.org/10.5281/zenodo.22675789\)](https://Research(doi.org/10.5281/zenodo.22675789))

Founder of Gnomon, delivering paid writing-assessment work for teachers and clients. Engineering experience spans data science, retrieval, multimodal training and evaluation.

EXPERIENCE

Gnomon | Founder & AI Systems Engineer | 2024 - Present

- Delivered paid teacher/client assessments; owned product design, issue triage, corrective fixes and final review.
- Designed rubric-compilation and evidence-verification components, grounding observations in source spans and preventing duplicate credit for repeated relations.
- Built exemplar retrieval for Gnomon and a separate LangGraph pipeline that retrieves compliance rules with SQLite FTS5.
- Engineered selection across 101,175 records and 47 source labels; prepared 10,888 training and 578 validation records using 44 data-quality checks.
- Led multimodal supervised fine-tuning with completion-only loss and chunked vocabulary projection; a separate 8B run processed 9,325 examples, including 2,134 with images.
- Designed agent scheduling with capacity and write-scope guardrails; directed React/TypeScript and FastAPI/SQLite delivery.

FPT Software | Data Science | 2025 - 2026

- Consolidated customer and business data, resolving missing records, inconsistencies and duplicates; analysed sales and operational data for trends, seasonality and anomalies.
- Compared forecasting models against baselines to recommend model selection.

FPT Education | Lecturer, AI for VFX | 2025 - 2026

- Delivered paid instruction on applying AI to visual-effects workflows.

FORWARD-DEPLOYED ENGINEERING - VFX AND AI | 2023 - 2025

OTSU Labs | Forward-Deployed 3D Generalist | 2025

- Embedded in studio production under Sparta VFX and Sparx supervision; owned render passes, asset management and look development, shipping shots to studio standard on deadline.

FPT, "Fragments of the Deep" | Creative Lead | 07/2024 - 2025

- Led the artist team on an Unreal Engine 5 animated short; optimised real-time models for a 30% performance gain.

Dihaan Media | Forward-Deployed 3D Artist | 2023

- Embedded with the client team on stage and booth environments (Sofitel Saigon Plaza), bringing new tools into their production workflow.

Contract engagements | VFX, education and government | 2023 - present

- AI, data and pipeline work with VFX and education client teams; contractor work on Vietnam's OCR book-to-e-book digitisation programme.

PROJECTS

failclosed-eval | Input validation | [View](#)

- Built fail-closed input validation for images, leakage and duplicates: 375 of 500 ELLIPSE-based payloads admitted, 125 malformed inputs refused.

policy-deck | Command classification | [View](#)

- Benchmarked 11 command-classification rules: 8 false positives/0 false negatives on 1,260 fitted cases; 2/0 on 190 separate cases. Classifier, not an execution sandbox.

relay-gate | Completion-claim detection | [View](#)

- Built a detector for false agent completion claims; benchmarked on MAST and AgentRewardBench, with 81 tests.

mcp-trajectory-judge | Tool-trajectory evaluation | [View](#)

- Replays synthetic MCP-style tool trajectories in a deterministic sandbox to check actions against completion claims.

Hackathon and competition entries | September 2026 | [View](#)

- Submitted a voice-agent entry with public code (uh-huh, 2 companion tools); scored an ARC White-Box Estimation leaderboard submission (arc-whest); entered Kaggle S6E9 with a scored public submission; OpenCV entry registered.

PUBLICATIONS & OPEN SOURCE

Constructing Language Judgments: Rubric Distillation and Evidence-Grounded Reasoning | [View](#)

EleutherAI lm-evaluation-harness | Submitted PR #4115 | [View](#)

- Added isolated DummyLM tests for construction, generation/error paths and chat templates; pull request open, not yet merged.

SKILLS

Evaluation: structured-output validation, pytest, GitHub Actions CI.

Retrieval and agents: LangGraph, theta-kNN exemplar retrieval, SQLite FTS5.

Machine Learning: Large Language Model (LLM) SFT, QLoRA, data-quality and leakage checks.

Languages, tools and serving: Python, PyTorch, PEFT/LoRA, TypeScript, React, FastAPI, SQLite.

Spoken languages: Vietnamese: native | English: professional proficiency | Japanese: intermediate.

EDUCATION

FPT University | Bachelor's degree.