

VO BA HOANG MINH

AI ENGINEER | AGENTS, EVALUATION & RETRIEVAL SYSTEMS

Ho Chi Minh City, Vietnam | Remote, worldwide | UTC+7

minhhoang250803@gmail.com | +84 902 391 936

github.com/anhminhzui-dev | [Portfolio \(anhminhzui-dev.github.io\)](https://Portfolio(anhminhzui-dev.github.io)) | [Research \(doi.org/10.5281/zenodo.22675789\)](https://Research(doi.org/10.5281/zenodo.22675789))

AI engineer who ships evaluation-first LLM systems: fail-closed evaluation harnesses, RAG retrieval and agentic orchestration, running in a paid product he founded and published as benchmarked open repositories with an open pull request to EleutherAI's lm-evaluation-harness. Data science at FPT Software; lecturer in AI applications at FPT Education.

EXPERIENCE

Gnomon | Founder & AI Systems Engineer | 2024 - Present

- Owned end-to-end delivery of paid writing-assessment outputs for teachers and clients — product design, issue triage and corrective fixes through final output review.
- Built the grading engine's production retrieval-augmented (RAG) layer — KD-tree theta-vector exemplar retrieval with a kNN-derived range constraint — and, in a second product, an agentic LangGraph pipeline with SQLite FTS5 (BM25-family) retrieval over a compliance-rule corpus.
- Engineered data selection across a 101,175-record inventory, 47 source labels; curated a 10,888/578 train/validation split via 44 registered checks for source, task, image, duplication and leakage risk.
- Led multimodal fine-tuning (SFT) with completion-only weighted loss, chunked vocabulary projection and accumulation-consistent normalization; an 8B fine-tune run processed 9,325 examples, 2,134 with images.
- Designed dependency-aware agent scheduling with capacity and write-scope guardrails; directed React/TypeScript and FastAPI/SQLite workflows from assignment through review and release.

FPT Software | Data Science | 2025 - 2026

- Prepared and consolidated historical customer and business data for analysis and predictive-model evaluation, addressing missing records, inconsistencies and duplication.
- Conducted exploratory analysis of sales, customer behaviour and operational data to identify trends, seasonality, relationships and unusual changes.
- Compared forecasting models against baselines to recommend model selection.

FPT Education | Lecturer, AI for VFX | 2025 - 2026

- Delivered paid instruction on applying AI to visual-effects workflows.

CONTRACT & CONSULTING WORK | alongside full-time roles

Forward-deployed engineering on client contracts | VFX and education industries

- Embedded with client teams to deliver AI, data and pipeline work on contract, across the VFX and education industries, alongside full-time roles.

Government digitisation programme | OCR conversion of printed books to e-books

- Worked with contractor teams on Vietnam's book-to-e-book OCR digitisation programme to preserve printed knowledge.

SELECTED ENGINEERING PROJECTS | 2026

LLM Evaluation & Data Quality | failclosed-eval | [View](#)

- Shipped an evaluation-input validator that fails closed on missing images, label leakage and duplicates: 375 of 500 ELLIPSE-based payloads admitted, 125 malformed inputs rejected.

Agent Command-Policy Evaluation | policy-deck | [View](#)

- Built an agent command-policy classifier, 11 rules scored like a model rather than run as an execution sandbox: 8 false positives and zero false negatives on a 1,260-case fitted deck, 2 and zero on a separate 190-case held-out deck.

Agent Evaluation | relay-gate | [View](#)

- Built an agent false-done detector benchmarked against the MAST and AgentRewardBench public agent-failure benchmarks, shipped with an 81-test suite.

COMPETITIONS | SEPTEMBER 2026

Voice-agent hackathon (lablab.ai) | ARC White-Box Estimation 2026 (Alcrowd) | Kaggle S6E9 | OpenCV | [View](#)

- Submitted a voice-agent entry with public code (uh-huh, 2 companion tools); scored an ARC white-box leaderboard submission (arc-whest); entered Kaggle S6E9 with a scored public submission; OpenCV entry registered.

RESEARCH & OPEN-SOURCE CONTRIBUTION

Constructing Language Judgments: Rubric Distillation and Evidence-Grounded Reasoning | [View](#)

- Authored a systems report connecting rubric compilation, evidence-verifying inference harnesses, verifier-filtered supervision and resource-aware multimodal adaptation.

EleutherAI lm-evaluation-harness | Submitted PR #4115 | [View](#)

- Added hermetic DummyLM tests for construction, generation/error paths and chat-template behavior; pull request open, not yet merged.

SKILLS

Evaluation: evaluation design, structured-output validation, test suites (pytest, GitHub Actions).

Retrieval and agents: LangGraph agent orchestration; theta-kNN exemplar retrieval; SQLite full-text/BM25-family retrieval; dependency-aware agent scheduling.

Training and data: QLoRA fine-tuning, multimodal SFT, data-quality and leakage checks.

Serving and MLOps: FastAPI, SQLite, GitHub Actions CI.

Languages and tools: Python, PyTorch, PEFT/LoRA, TypeScript, React.

EDUCATION

FPT University | Bachelor's degree.

LANGUAGES

Vietnamese: native | English: professional proficiency | Japanese: intermediate.