

VO BA HOANG MINH

AI ENGINEER | MODEL INTEGRATION & EVALUATION

Ho Chi Minh City, Vietnam | Remote + onsite (Vietnam) | UTC+7

Portfolio: github.com/anhminhzui-dev | minhhoang250803@gmail.com | +84 902 391 936

AI systems engineer specializing in rubric distillation, evidence-grounded reasoning, and training/inference pipelines. Lead Gnomon's architecture, implementation direction and acceptance testing, with seven public repositories spanning evaluation, agent tooling and model operations.

TECHNICAL SKILLS

Software: Python, SQL/SQLite, TypeScript, FastAPI, REST APIs, React.

ML: PyTorch, PEFT/QLoRA, SFT, data QA, structured-output validation.

Testing & automation: pytest, GitHub Actions.

ENGINEERING EXPERIENCE

Gnomon | IELTS Rubric Distillation & Reasoning Systems

Founder & AI Systems Engineer | 2024-present

- **Rubric distillation:** Synthesized IELTS writing criteria into a compiled, versioned rubric bank shared by prompts and scoring, linking evidence, judgments and teacher review.
- **Reasoning engineering:** Built exact source-span validation and relation-level evidence accounting to prevent duplicate credit and distinguish unmeasured evidence from measured failure.
- **Data engineering:** Designed selection over a 101,175-record inventory across 47 source labels, with task/image checks and held-out exclusions. Prepared 10,888 training and 578 validation records with duplicate/leakage checks.
- **Training systems:** Engineered multimodal SFT with completion-masked weighted loss, chunked frozen-head vocabulary projection and microbatch-consistent loss normalization.
- **Judgment distillation:** Designed verifier-guided training-target selection from model attempts and reviewed repairs; corrections require improved verifier reward and matching first-pass examples.
- **Inference engineering:** Developed low-rank adapter injection and inference-time merging for the earlier ModernBERT path, alongside DirectML-compatible AdamW training.

Research: Beyond Score Prediction: Evidence-First Reasoning for Writing Assessment

Independent technical report (2026): rubric compilation, reasoning, data and training systems. [Paper](#)

SELECTED RECENT PROJECTS (3 OF 7)

Evaluation Data Integrity | Python | 2026 | [Code](#)

- Validated image inputs and label leakage. Controlled input-validation test: 500 ELLIPSE-based payloads; 375 valid admitted and 125 deliberately malformed refused.

Agent Guardrail Evaluation | Python | 2026 | [Code](#)

- Measured command-policy errors on pinned synthetic decks: 8 false positives / 0 false negatives on 1,260 fitted rows; 2 / 0 on 190 separate, partly rule-aware rows. Classifier, not an execution sandbox.

Model Release Validation | Python | 2026 | [Code](#)

- Implemented candidate-versus-baseline GO/HOLD checks for criterion regressions, sample size and coverage; preserves input hashes and exercises decisions with synthetic failure fixtures.

EDUCATION & LANGUAGES

FPT University | Bachelor of Arts | Ho Chi Minh City, Vietnam.

Languages: Vietnamese (native), English (professional proficiency), Japanese (intermediate).