

# VO BA HOANG MINH

## 3D / VFX ARTIST | AI WORKFLOW ENGINEER | REMOTE, UTC+7

Ho Chi Minh City, Vietnam | Remote, UTC+7

minhhoang250803@gmail.com | +84 902 391 936

[github.com/anhminhzui-dev](https://github.com/anhminhzui-dev) | [Portfolio \(anhminhzui-dev.github.io\)](https://portfolio.anhminhzui-dev.github.io) | [Research \(doi.org/10.5281/zenodo.22675789\)](https://doi.org/10.5281/zenodo.22675789)

3D/VFX production experience in asset preparation, look development, render passes and studio delivery. I also teach AI for VFX and build Python-based AI evaluation and review workflows.

## 3D & VFX PRODUCTION | 2023 - 2025

### OTSU Labs | 3D Generalist | 2025

- Worked in VFX production under Sparta VFX and Sparx supervision; handled render passes, asset management and look development under studio deadlines; converted realistic assets into stylised anime visuals with the 2D team.

### FPT, "Fragments of the Deep" | Creative Lead | 07/2024 - 2025

- Led the artist team on an Unreal Engine 5 animated short, owning storyboard and camera direction; optimised real-time models.

### Dihaan Media | 3D Artist | 2023

- Embedded with the client production team on stage and booth environments for Sofitel Saigon Plaza; introduced IES-based lighting and new tools into their workflow.

## Contract engagements | VFX, education and government | 2023 - present

- AI, data and pipeline work with VFX and education client teams; contractor work on Vietnam's OCR book-to-e-book digitisation programme.

## EXPERIENCE

### Gnomon | Founder & AI Systems Engineer | 2024 - Present

- Delivered paid teacher and client assessments; owned product design, triage, fixes and final review.
- Designed the facts-not-grades architecture: the model emits evidence grounded in source spans, a deterministic scorer maps facts to the result, and an abstain layer refuses on thin evidence.
- Built the fail-closed evaluation harness: 2,232 criterion units across 629 essays, 1,866 admitted and 366 refused before scoring; LLM-as-judge runs on a frozen, hash-checked evidence window.
- Built KD-tree exemplar retrieval for Gnomon and a separate LangGraph pipeline that retrieves compliance rules with SQLite FTS5.
- Engineered selection across 101,175 records and 47 source labels; prepared 10,888 training and 578 validation records using 44 data-quality checks.
- Led multimodal supervised fine-tuning: QLoRA on a 27B open-weight model (4-bit NF4) with completion-only loss and chunked vocabulary projection; a separate 8B run processed 9,325 examples, 2,134 with images.
- Directed a two-provider coding-agent fleet under policy-as-code guardrails (resource-aware admission, spend/irreversibility classifier scored on a hash-pinned 1,219-case deck, output-hash delivery checks); shipped the React/TypeScript and FastAPI/SQLite product: 43 router modules, 79 migrations, TOTP MFA, 1,117 backend tests.
- Presentation design and slide reconstruction of a ten-slide Gnomon sales deck: action titles, data visualization, rationale log.

### FPT Software | Data Science | 2025 - 2026

- Consolidated customer and business data (missing records, inconsistencies, duplicates); analysed sales and operational data for trends, seasonality and anomalies; compared forecasting models to recommend selection.

### FPT Education | Lecturer, AI for VFX | 2025 - 2026

- Delivered paid instruction on applying AI to visual-effects workflows.

## PROJECTS

### failclosed-eval | Input validation | [View](#)

- Built fail-closed input validation for images, leakage and duplicates: 375 of 500 ELLIPSE-based payloads admitted, 125 malformed inputs refused; 239 tests, GitHub CI.

### policy-deck | Command classification | [View](#)

- Benchmarked 11 command-classification rules: 8 false positives/0 false negatives on 1,260 fitted cases; 2/0 on 190 separate cases; 455 tests. Classifier, not an execution sandbox.

### Hackathon and competition entries | September 2026 | [View](#)

- Submitted a voice-agent entry with public code (uh-huh, 2 companion tools); scored an ARC White-Box Estimation leaderboard submission (arc-whest); Kaggle S6E9 public score 0.94151; OpenCV entry registered.

## PUBLICATIONS & OPEN SOURCE

### Constructing Language Judgments: Rubric Distillation and Evidence-Grounded Reasoning | [View](#)

### EleutherAI lm-evaluation-harness | Submitted PR #4115 | [View](#)

- Added isolated DummyLM tests for construction, generation/error paths and chat templates.

## SKILLS

Evaluation: LLM-as-judge on hashed evidence, structured-output validation, pytest, GitHub Actions CI.

Retrieval and agents: LangGraph, KD-tree theta-kNN exemplar retrieval, SQLite FTS5, policy-as-code agent guardrails.

Machine Learning: Large Language Model (LLM) SFT, QLoRA, data-quality and leakage checks.

3D and VFX: Blender, Maya, Houdini, Unreal Engine 5, asset management, look development, lighting, render passes.

Engineering: Python, TypeScript, PyTorch, PEFT/LoRA, React, FastAPI, SQLite.

Spoken languages: Vietnamese: native | English: professional proficiency | Japanese: intermediate.

## EDUCATION

FPT University | Bachelor's degree.